

### Q1(a)(i) Differences between TensorFlow and PyTorch

#### Community and Ecosystem

While *TensorFlow* has a strong ecosystem backed by Google, *PyTorch* is loved by the research community, is widely used in academia and cutting-edge AI research, and is backed by Meta.

#### Deployment

*PyTorch* initially focused on research, but with TorchServe and ONNX support, deployment has become much easier while *TensorFlow* is better suited for production and deployment, especially with TensorFlow Serving, TensorFlow Lite and TensorFlow.js.

#### Computation Graphs

*PyTorch* uses dynamic computation graphs (define-and-run), you create the graph on the fly as operations are executed while *TensorFlow* uses static computation graphs (define-before-run), you define the model graph first then run it.

#### Debugging

While *PyTorch* is easier to debug, intuitive and flexible – making it ideal for research *TensorFlow* is less flexible for debugging and experimentation.

### Q1(a)(i)-(b) When to choose TensorFlow and PyTorch

- ❖ When focusing on **deployment and scalability**, go with *TensorFlow*.
- ❖ When the focus is on **research, flexibility, debugging ease, and fast experimentation**, choose *PyTorch*.

### **Q1(a)(ii)** *Use cases for Jupyter Notebooks in AI Development*

#### **Prototyping and Testing AI Models**

With Jupyter Notebooks, developers can quickly build, test and iterate on models as it is:

- Ideal for experimentation with different neural hyperparameters
- Combine code, outputs and documentation in one place for reproducibility.

#### **Data Exploration and Processing**

Data Scientists use Jupyter Notebooks to load, explore, and clean data before training AI Models.

With Jupyter Notebooks one can:

- Visualize data using libraries like Seaborn or Matplotlib.
- Perform data wrangling, feature selection and do analysis step by step.
- Tweak code as needed to get the desired output.

**Q1(a)(iii)** *How spaCy enhances NLP tasks compared to basic Python string operations.*

### **Advanced NLP Features.**

Unlike basic Python string operations, spaCy has advanced Natural Language Processing features that allow:

- Tokenization – While basic Python allows manual splitting spaCy has smart tokenization that handles abbreviations, punctuations, etc.
- Lemmatization – Converts words to their base.
- Named Entity Recognition (NER) – Detects people, dates, locations, etc.
- Dependency Parsing – Understands grammatical relationships.
- Part of Speech Tagging – Identifies verbs, adjectives, nouns, etc.

### **Linguistic Understanding**

Basic Python works at the character or word level, but it does not understand language structure while spaCy performs linguistic analysis, recognizing parts of speech, the syntax and entities.

#### **Example:**

*“A tasty burger.”*

Basic python will analyze it as “A “,”tasty“, “burger. “(no grammar or meaning while spaCy will identify “A” as article a, “tasty” as adjective, “burger” as a noun and their grammatical relationship.

### **Speed and Efficiency**

While in Python string methods one would have to manually loop over text, spaCy is highly optimized in Cython, making it faster and more memory efficient.

**Q1(b)** In terms of Target applications, Ease of use for beginners and community support, compare Scikit-learn and TensorFlow.

Indicator	Scikit-learn	TensorFlow
Target Applications	Designed for <b>classical machine learning</b> algos like linear regression, random forests, clustering, decision trees, and SVMs.	Built for <b>deep learning and neural networks</b> , supporting large scale models like RNNs, CNNs and Transformers.
Ease of Use for Beginners	<b>Very beginner friendly</b> . It has simple, consistent API and models can be trained in just a few lines of code.	<b>Moderate learning curve</b> - more complex concepts (tensors, graphs, sessions), though Keras (now part of TensorFlow) simplifies it significantly.
Community Support	It has a <b>large and active data science community</b> making it excellent for documentation and tutorials for traditional machine learning.	<b>Massive global community</b> backed by Google therefore making it an extensive ecosystem for deep learning, deployment, and production use.

## Q2(a)

Classical ML with Scikit-learn visualization output:



## Q2(b)

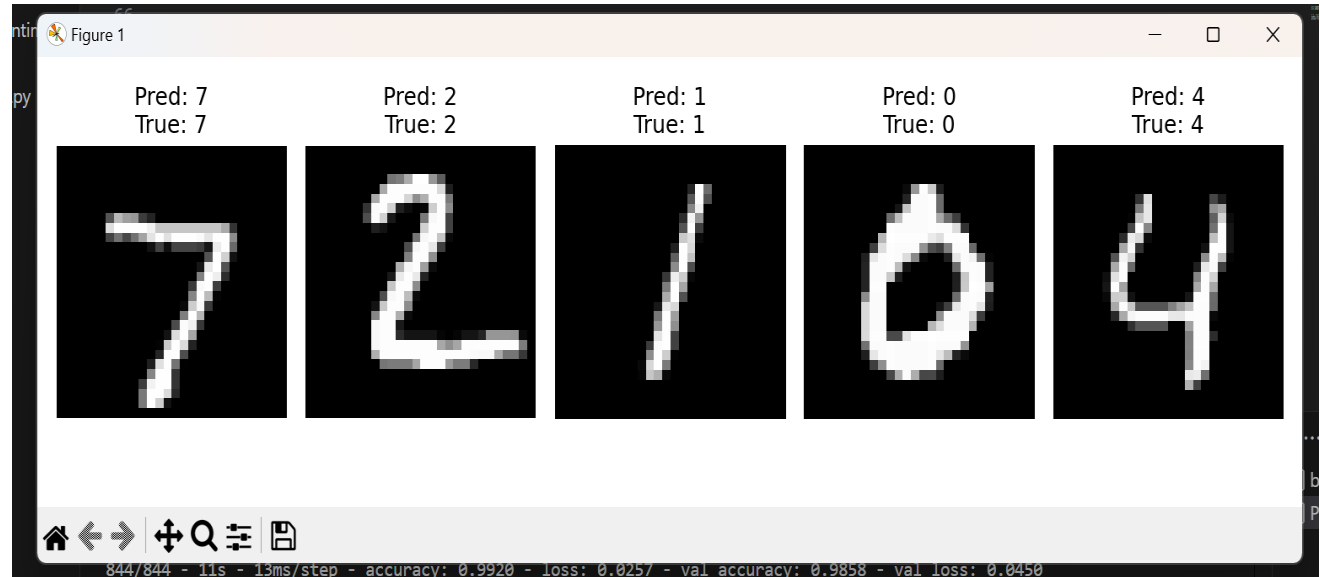
*Training CNN model accuracy results:*

```
Training the CNN model...

Epoch 1/5
844/844 - 10s - 12ms/step - accuracy: 0.9414 - loss: 0.1902 - val_accuracy: 0.9862 - val_loss: 0.0478
Epoch 2/5
844/844 - 9s - 11ms/step - accuracy: 0.9821 - loss: 0.0581 - val_accuracy: 0.9855 - val_loss: 0.0453
Epoch 3/5
844/844 - 11s - 13ms/step - accuracy: 0.9877 - loss: 0.0400 - val_accuracy: 0.9905 - val_loss: 0.0330
Epoch 4/5
844/844 - 11s - 13ms/step - accuracy: 0.9904 - loss: 0.0315 - val_accuracy: 0.9883 - val_loss: 0.0402
Epoch 5/5
844/844 - 11s - 13ms/step - accuracy: 0.9920 - loss: 0.0257 - val_accuracy: 0.9858 - val_loss: 0.0450

✅ Test Accuracy: 98.80%
1/1  0s 115ms/step
```

*Deep Learning with TensorFlow output:*



## Q2(c)

NLP with spaCy output:

```
39
40 # 4. Process each review

PROBLEMS 1 OUTPUT DEBUG CONSOLE TERMINAL PORTS

/Task3/nlp_spacy_sentiment.py
📄 Review: I absolutely love my new Apple iPhone! The camera quality is amazing.
💎 Named Entities: [('Apple', 'ORG')]
💬 Sentiment: Positive 😊
-----
📄 Review: The Samsung Galaxy screen cracked within a week. Totally disappointed.
💎 Named Entities: [('Samsung Galaxy', 'ORG'), ('a week', 'DATE')]
💬 Sentiment: Negative 😞
-----
📄 Review: This Logitech mouse is super smooth and comfortable to use.
💎 Named Entities: [('Logitech', 'ORG')]
💬 Sentiment: Positive 😊
-----
📄 Review: I hate the battery life on my Dell laptop. It barely lasts 2 hours!
💎 Named Entities: [('Dell', 'ORG'), ('2 hours', 'TIME')]
💬 Sentiment: Negative 😞
-----
📄 Review: The Sony headphones have great sound but the build quality feels cheap.
💎 Named Entities: [('Sony', 'ORG')]
💬 Sentiment: Neutral 😐
-----
```

**Q3(a)** *Potential biases in MNIST or Amazon Review model and how TensorFlow or SpaCy can mitigate these biases.*

### **Ethical Considerations**

*Potential Biases.*

#### **Amazon Reviews (spaCy NLP)**

- *Bias Source:* Rule-based sentiment models depend on specific keywording.

Example

“hate” = negative, “love” = positive.

- *Risk:* This may misclassify sarcasm, gendered languages, or cultural phrases.
- *Impact:* This could, in addition, unfairly flag or recommend products if reviews use mixed sentiment or dialect-specific language.

*Mitigation Tool*

#### **spaCy Rule-based Systems**

One can use custom lexicons, dependency rules, or contextual filters to avoid one-size-fits-all sentiment scoring.

Example

Adjust sentiment weights based on negations (“not good”) or sarcasm cues (“yeah right”).



**Q3(b) Buggy Code:** A provided TensorFlow script has errors (e.g., dimension mismatches, incorrect loss functions). Debug and fix the code.

### Troubleshooting challenge:

```
19 model = tf.keras.Sequential([
20     tf.keras.layers.Dense(64, input_shape=(28, 28), activation='relu'),
21     tf.keras.layers.Dense(10, activation='softmax')
22 ])
23
24 model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
25 model.fit(x_train, y_train, epochs=5)
```

### Issues:

- Input shape mismatch- The input should be (28,28,1) or (784) flattened, not (28,28)
- Loss mismatch- Using categorical\_crossentropy but y-train likely contains integer labels instead of one-hot vectors.

```
import tensorflow as tf
from tensorflow.keras import layers, models
# Fix: Flatten input properly for Dense layers
model = models.Sequential([
    layers.Flatten(input_shape=(28, 28)), # Correct input shape
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])
# Fix: Use correct loss for integer labels
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
# Example dataset
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
# Normalize
x_train, x_test = x_train / 255.0, x_test / 255.0
# Train
model.fit(x_train, y_train, epochs=5, validation_split=0.1)
```

### Fixed Version

- Flattened input (Flatten() layer)
- Correct loss function (sparse\_categorical\_crossentropy)
- Normalized data for stable training.